

Data Persistence

By
Raúl Mazo
Paola Vallejo


Data persistence

- Persistence is the ability to **save the information** of a program for **reuse** at a **later** time. Persistence is the **permanence** of information
- User:
 - Save, Save as, Recover ...
- Programmer:
 - mechanism responsible for the **backup** and **recovery** of data
 - a program can be terminated **without losing its data** and execution status
 - for a data store to be considered persistent, it must write to **non-volatile storage**
 - **serialization** of the data and the reverse process of retrieving the data from the serialized information


Why data persistence?

- For later use
- To maintain the status
- For logging purposes
- To further process and derive knowledge
- Data can be stored, read, updated/modified, and deleted


Ways of data persistent



runtime memory
cache memory



file



database

Data storage formats

Plain-text, XML, JSON, tables, text files, images

Ways of data persistent

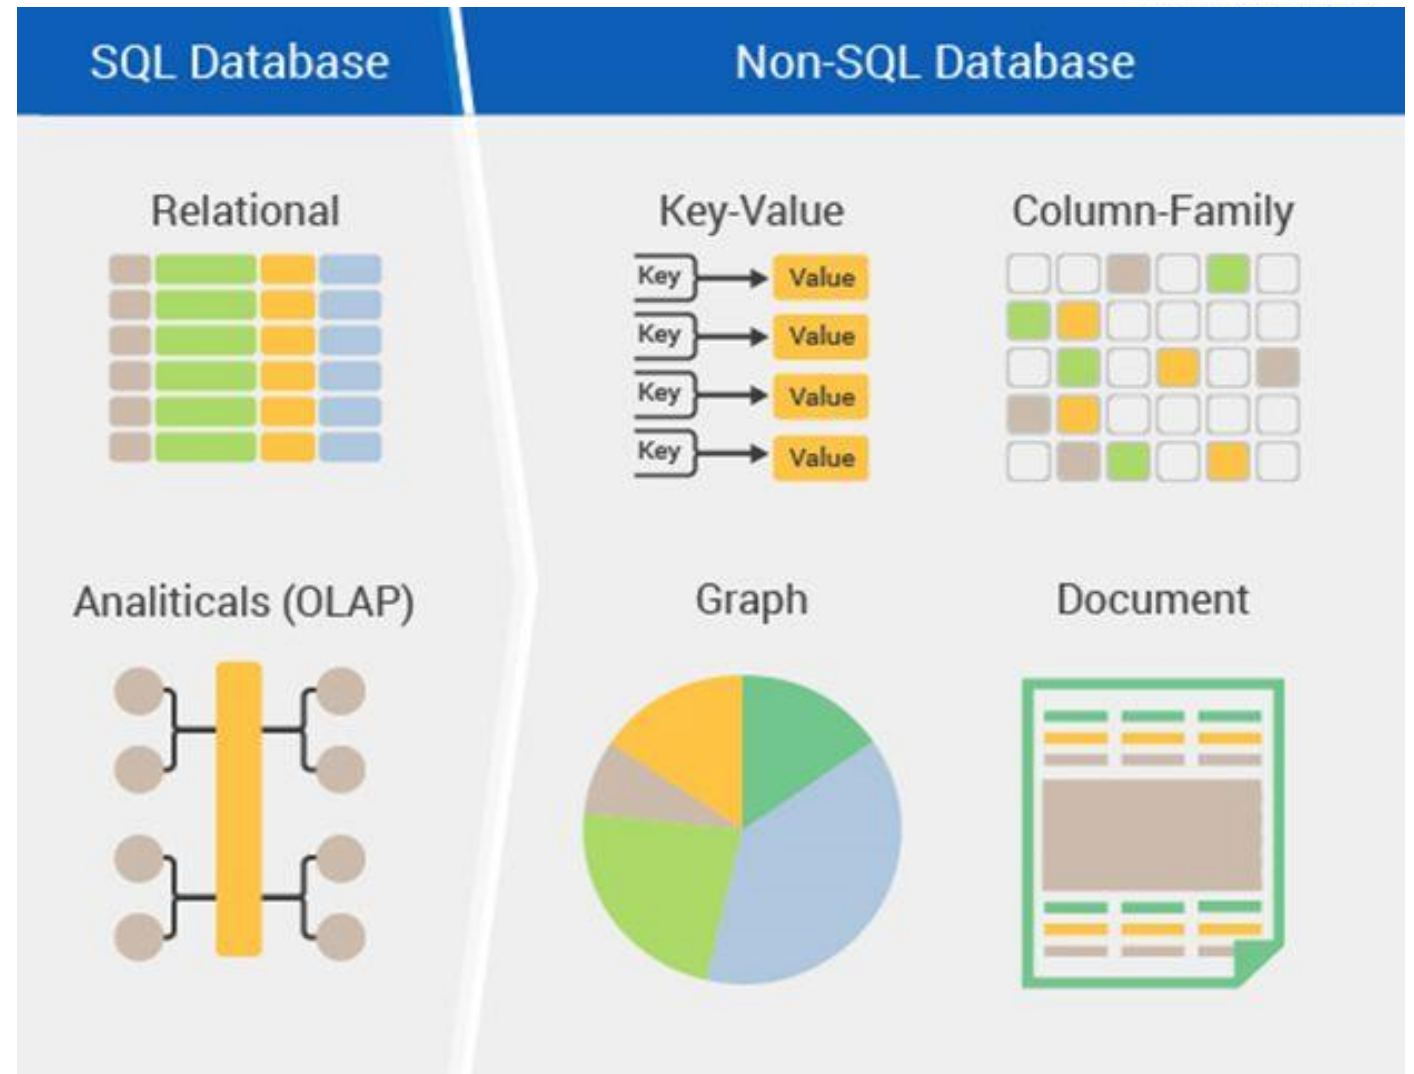
- **Pure in-memory**, such as cache memory
 - High speed and availability
 - Few data
 - No persistence at all
- **In-memory with periodic snapshots**, such as Redis
 - Periodic snapshots to disk at a configurable interval
- **Disk-based** with update-in-place writes, such as MySQL or MongoDB
- **Commitlog-based**, such as all traditional OLTP databases (Oracle, SQL Server ...)

Database type

- Object-oriented
- Network
- Hierarchical

SQL

Non-SQL



Contexts

NoSQL



Gaming



Social



IoT



Web



Mobile



Enterprise




Graph



Key/value store



Document
database



Column family store

SQL



Web




Mobile




Enterprise



Data mart




Relational table storage



Relationships use joins

Database Server



- **Back-end system** of a database application using client/server architecture.
- Performs tasks such as data analysis, storage, data manipulation, and other non-user specific tasks.

Database Management System



- Software for creating and managing databases.
- Interface between the database and end users or application programs. Ex. MySQL, PostgreSQL, Microsoft Access, SQL Server, Oracle.

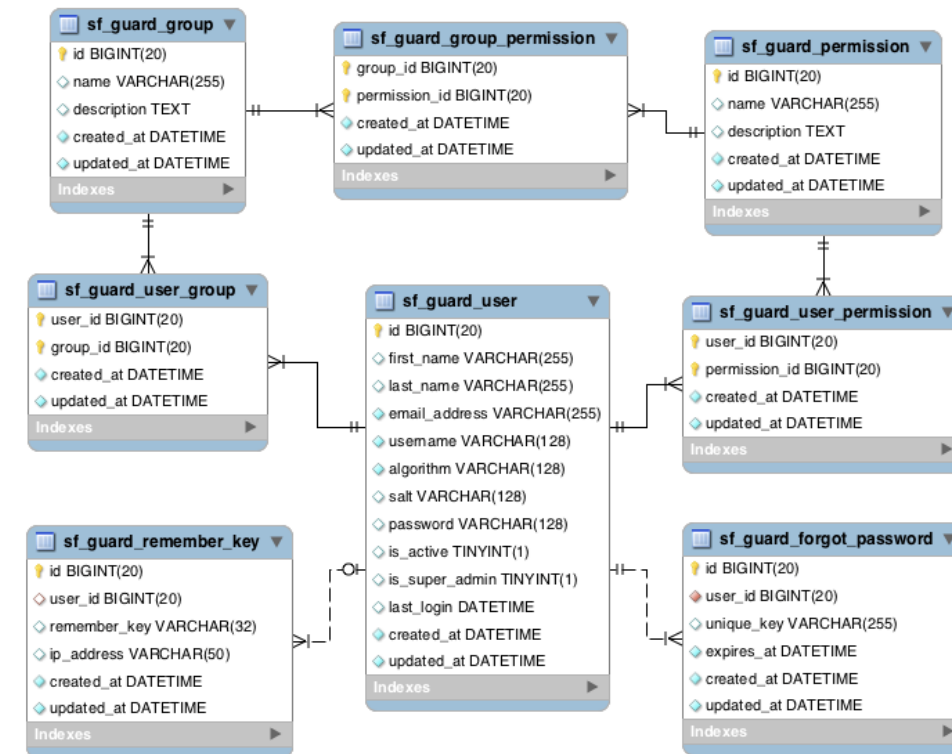
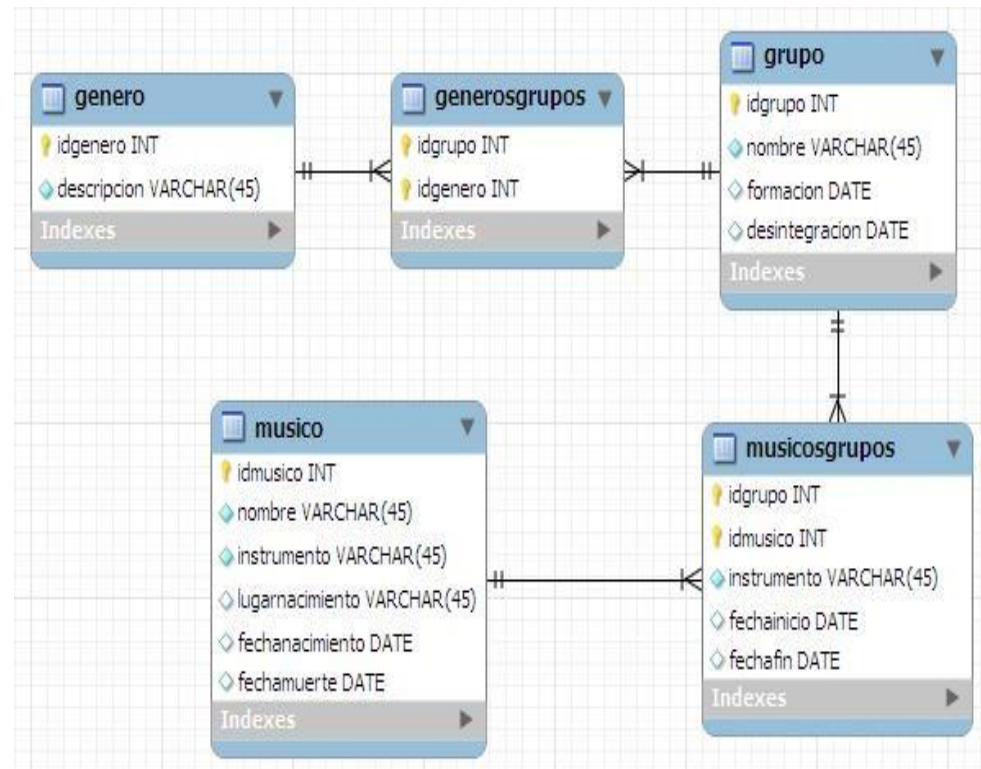
Relational models

General concept

The relational data model organizes and represents data in the form of tables and relationships:

Logical representation	Physical representation	Relational model
Table	Sequential File	Relation
Row	Register	Tuple
Column	Field	Attribute


Example - Relational model



Non-SQL models

SQL database issues

Major problems that make it ineffective



Huge growth of information on the Internet. The growth of data from 2007 to 2010 is almost 25 times.

Connectivity. Over time, information is becoming more and more connected.

Semi-structure. Information that has some mandatory attributes, but many optional attributes. As the information grew, there was a need to increase the columns of the table, which would lead to sparse tables.

Response to relational inefficiency

❓ NoSQL databases that were initially created in response to the needs for better scalability, lower latency and greater flexibility in the era of bigdata and cloud computing.

These non-functional aspects are the main reason for the use of NoSQL databases.

NoSQL databases are structures that allow storing information in situations where relational databases generate scalability and performance problems when thousands of concurrent users and millions of daily queries occur.

NoSQL advantages

Run on machines with few resources:	<ul style="list-style-type: none"> • Require little computation, so they can be mounted on lower cost machines.
Horizontal scalability:	<ul style="list-style-type: none"> • To improve performance, more nodes are added, with the only operation being to indicate to the system which nodes are available.
Handle large amounts of data:	<ul style="list-style-type: none"> • Distributed structure, in many cases by means of Hash tables.
Does not generate bottlenecks:	<ul style="list-style-type: none"> • SQL systems need to transcribe each statement in order to be executed, which is a common entry point, which can slow down the system when faced with many requests.

Non-SQL


Do not use SQL as a query language

- Most NoSQL databases avoid using this type of language or use it as a support language. To give some examples, Cassandra uses the CQL language, MongoDB uses JSON or BigTable uses GQL.


Do not use fixed structures such as tables for data storage.

- They allow the use of other types of information storage models such as key-value systems, objects or graphs.


INNER JOIN




FULL JOIN



LEFT JOIN



RIGHT JOIN



JOIN
operations
are not
usually
allowed



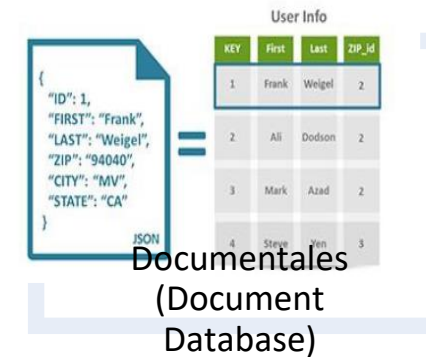

- When having such an extremely large volume of data it is often desirable to avoid JOINS.
- This is because, when the operation is not a key lookup, the overhead can become very costly.
- The most straightforward solutions are to denormalize the data, or to perform the JOIN in software, at the application layer.

Distributed
architecture

- Relational databases are usually centralized in a single machine or in a master-slave structure, however in NoSQL cases the information may be shared in several machines by means of distributed Hash table mechanisms.

NoSQL Overview

NoSQL database types



NoSQL database examples

Key Value	<ul style="list-style-type: none"> • Berkeley DB, Tokyo Tyrant, Voldemart, Crassandra.
Big Table	<ul style="list-style-type: none"> • Google BigTable, HBase
Document	<ul style="list-style-type: none"> • Mongo DB, Couch DB.
Graph	<ul style="list-style-type: none"> • Neo4j, InfoGrid o Virtuoso

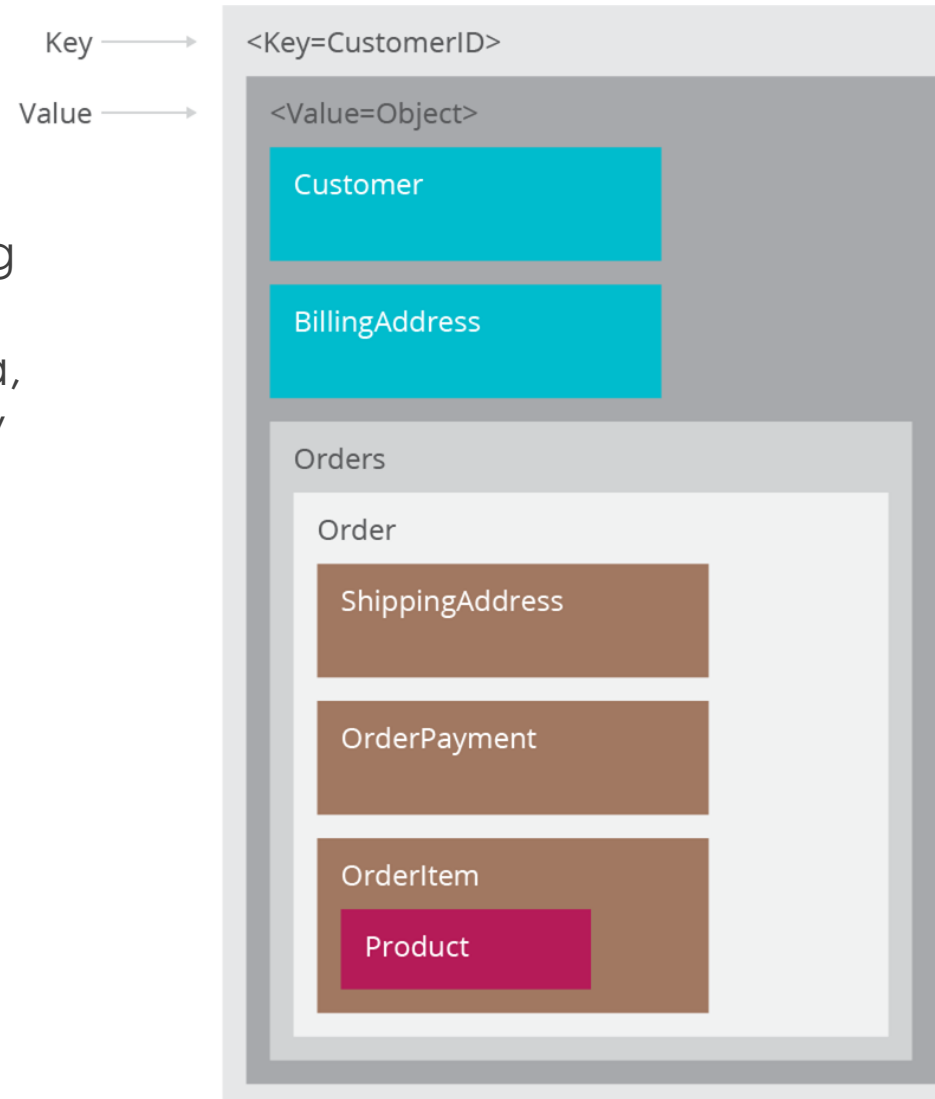
Key - Value

- They store tuples containing a key and its value. When you want to retrieve a data, you simply search for its key and retrieve the value.

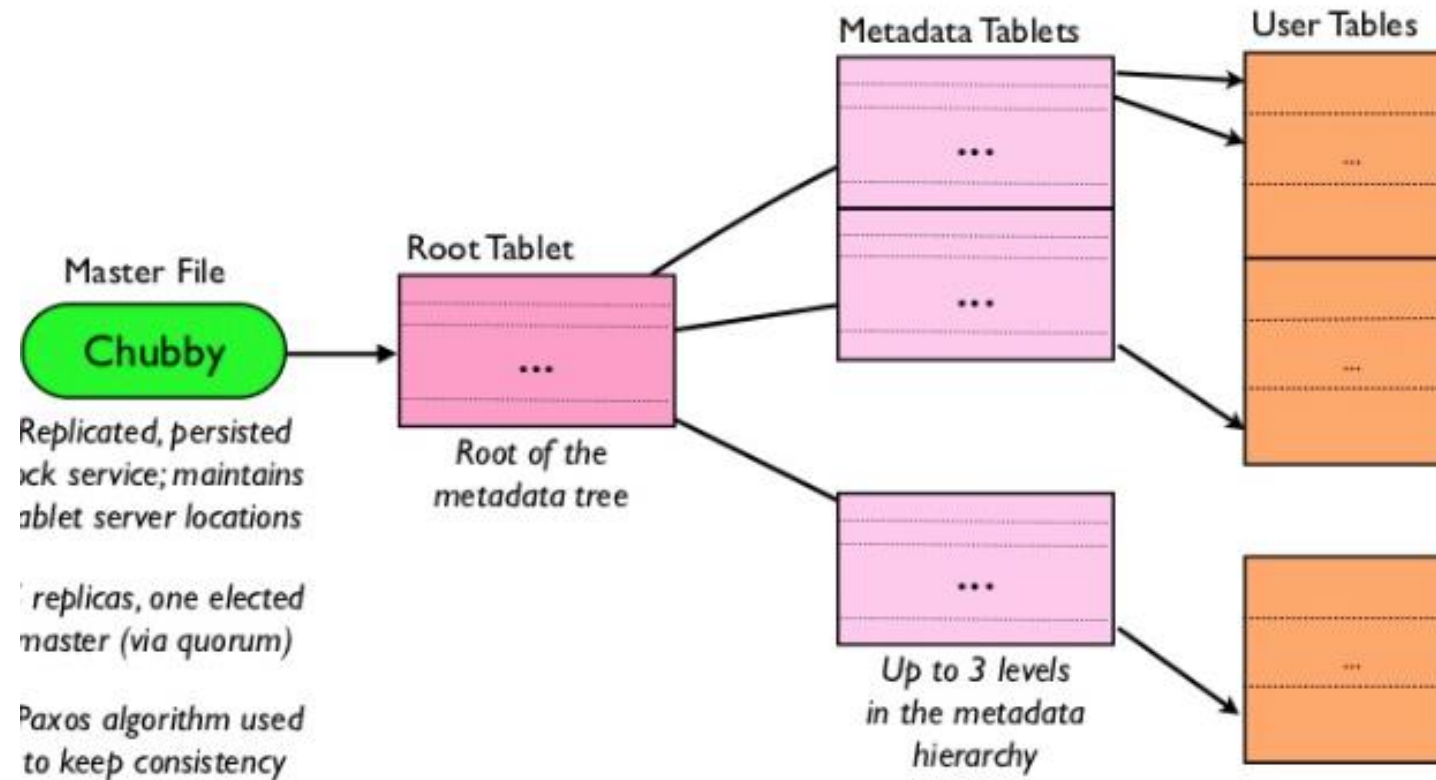
- Hash tables

- DynamoDB


- Redis




Bigtable



Documents



reference relations



embedded documents

```

{
  "firstName": "Shane",
  "lastName": "Johnson",
  "skills": ["Big Data", "Java",
    "NoSQL"],
  "experience": [
    {
      "role": "Technical Marketing",
      "company": "Red Hat"
    },
    {
      "role": "Product Marketing",
      "company": "Couchbase"
    }
  ]
}

```

? Semi-structured data, i.e. documents. These data are stored in some standard format such as XML, JSON or BSON.

? MongoDB

? CouchDB

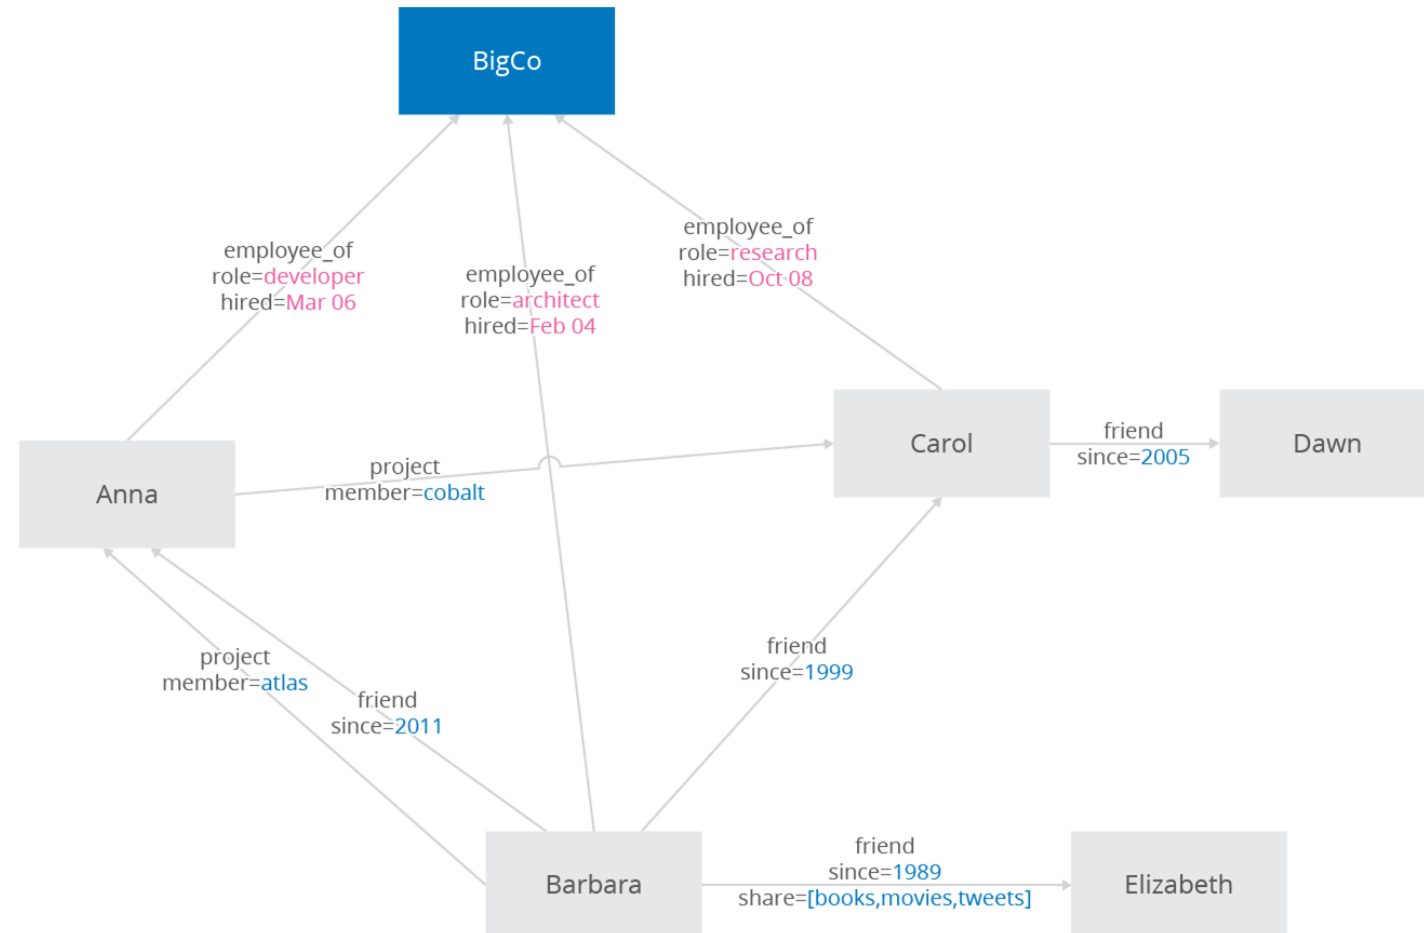
Graph

- Based on graph theory, they use nodes and edges to represent the stored data. They are very useful for storing information in models with many relationships, such as networks and social connections.

- Queries using indexes.

- Infinite Graph

- Neo4j



When to use NoSQL DB

Some of the reasons that may lead us to use NoSQL databases instead of the classic SQL databases are:

The volume of data grows very rapidly at specific moments, and can exceed one terabyte of information.

The scalability of the relational solution is not feasible both at a cost and technical level.

High peak usage of the system by users on multiple occasions.

The database schema is not homogeneous, i.e. when in each data insertion the information stored may have different fields.

Example 1

What would be the structure of a JSON document?

SCHEMA:

Column	Format
first_name	VARCHAR
last_name	VARCHAR
birth_date	DATE

TABLE / ROW:

first_name	last_name	birth_date
Shane	Johnson	08/12/1979

```
{  
  "firstName": "Shane",  
  "lastName": "Johnson",  
  "birthDate": "08/12/1979"  
}
```

Example 2



Los documentos JSON de ejemplo que representan a estas personas son, para el caso de la persona:

```
{
  nombre: "V́ctor Cuervo",
  edad: 38
}
```

Y para el caso del domicilio es:

```
{
  calle: "Alcala, 15",
  codigo: 28022,
  ciudad: "Madrid"
}
```

```
{
  nombre: "V́ctor Cuervo",
  edad: 38,
  direcci3n: {
    calle: "Alcala, 15",
    codigo: 28022,
    ciudad: "Madrid"
  }
}
```

```
db.personas.find({nombre:"V́ctor Cuervo"},{direccion:1})
```

```
{
  _id: 1,
  nombre: "V́ctor Cuervo",
  edad: 38
}
```

Y ese id ser1 utilizado dentro del documento del domicilio:

```
{
  userid: 1,
  calle: "Alcala, 15",
  codigo: 28022,
  ciudad: "Madrid"
}
```

```
var id = db.personas.find({nombre:"V́ctor Cuervo"},{_id:1})
db.domicilios.find({userid:id})
```

Exercise



En este caso los documentos JSON con los que contamos serán, por un lado la entrada del blog:

```
{
  title: "Línea de Código",
  url: "http://lineadecodigo.com",
  text: "Aprende a Programar"
}
```

Y por otro los N comentarios que existan:

```
{
  name: "Carlos Camacho",
  created_on: ISODate("2015-12-01T10:01:22Z"),
  comment: "Me gusta tu blog"
}

{
  name: "Fran Honrubia",
  created_on: ISODate("2015-12-01T14:15:10Z"),
  comment: "Gran trabajo"
}
```

```
{
  title: "Línea de Código",
  url: "http://lineadecodigo.com",
  text: "Aprende a Programar",
  comments: [
    {
      name: "Carlos Camacho",
      created_on: ISODate("2015-12-01T10:01:22Z"),
      comment: "Me gusta tu blog"
    },
    {
      name: "Fran Honrubia",
      created_on: ISODate("2015-12-01T14:15:10Z"),
      comment: "Gran trabajo"
    }
  ]
}
```

```
db.post.find({title:"Línea de Código"},{comments:1});
```

```
{
  _id:1,
  title: "Línea de Código",
  url: "http://lineadecodigo.com",
  text: "Aprende a Programar"
}
```

Y por otro lado cada uno de los comentarios con el _id como foreign key.

```
{
  blog_entry: 1,
  name: "Carlos Camacho",
  created_on: ISODate("2015-12-01T10:01:22Z"),
  comment: "Me gusta tu blog"
}

{
  blog_entry: 1,
  name: "Fran Honrubia",
  created_on: ISODate("2015-12-01T14:15:10Z"),
  comment: "Gran trabajo"
}
```

```
var post_id = db.post.find({title:"Línea de Código"},{_id:1});
db.comments.find({blog_entry: post_id}).foreach(doc) {
  print (doc.name + doc.comment)
}
```

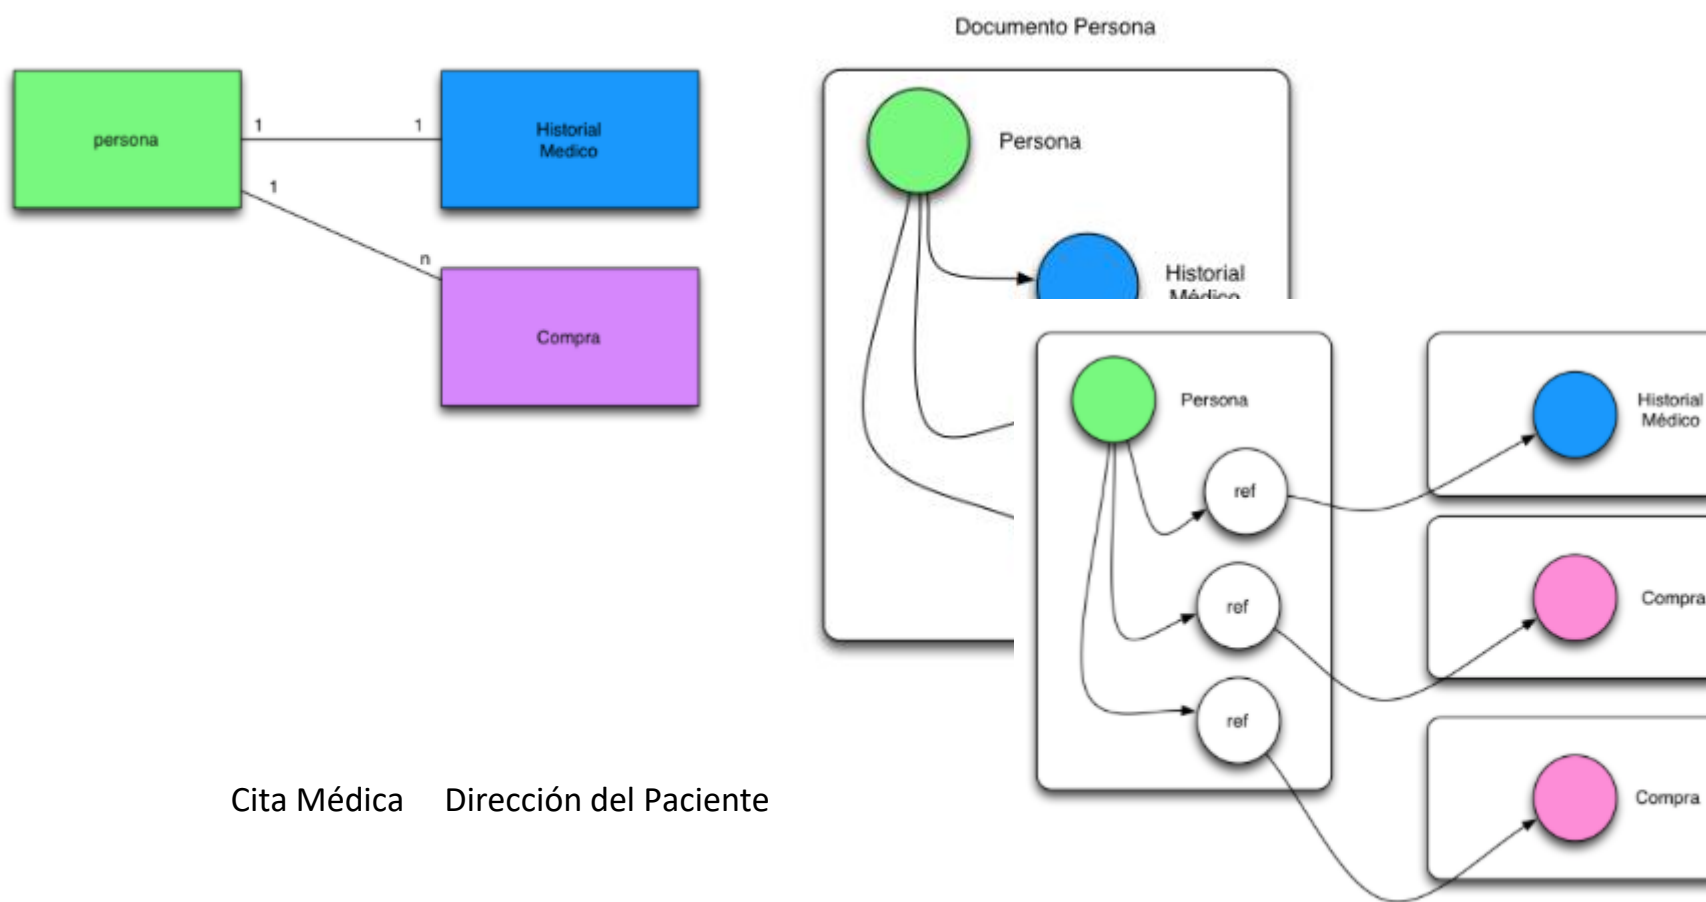
Example 3

USERS		
<i>ID</i>	<i>First</i>	<i>Last</i>
1	Shane	Johnson

USER SKILLS	
<i>User ID</i>	<i>Skill Name</i>
1	Big Data
1	Java
1	NoSQL

USER EXPERIENCE		
<i>User ID</i>	<i>Role</i>	<i>Company</i>
1	Technical Mktg	Red Hat
1	Product Mktg	Couchbase

```
{
  "firstName": "Shane",
  "lastName": "Johnson",
  "skills": ["Big Data", "Java",
    "NoSQL"],
  "experience": [
    {
      "role": "Technical Marketing",
      "company": "Red Hat"
    },
    {
      "role": "Product Marketing",
      "company": "Couchbase"
    }
  ]
}
```



Technologies



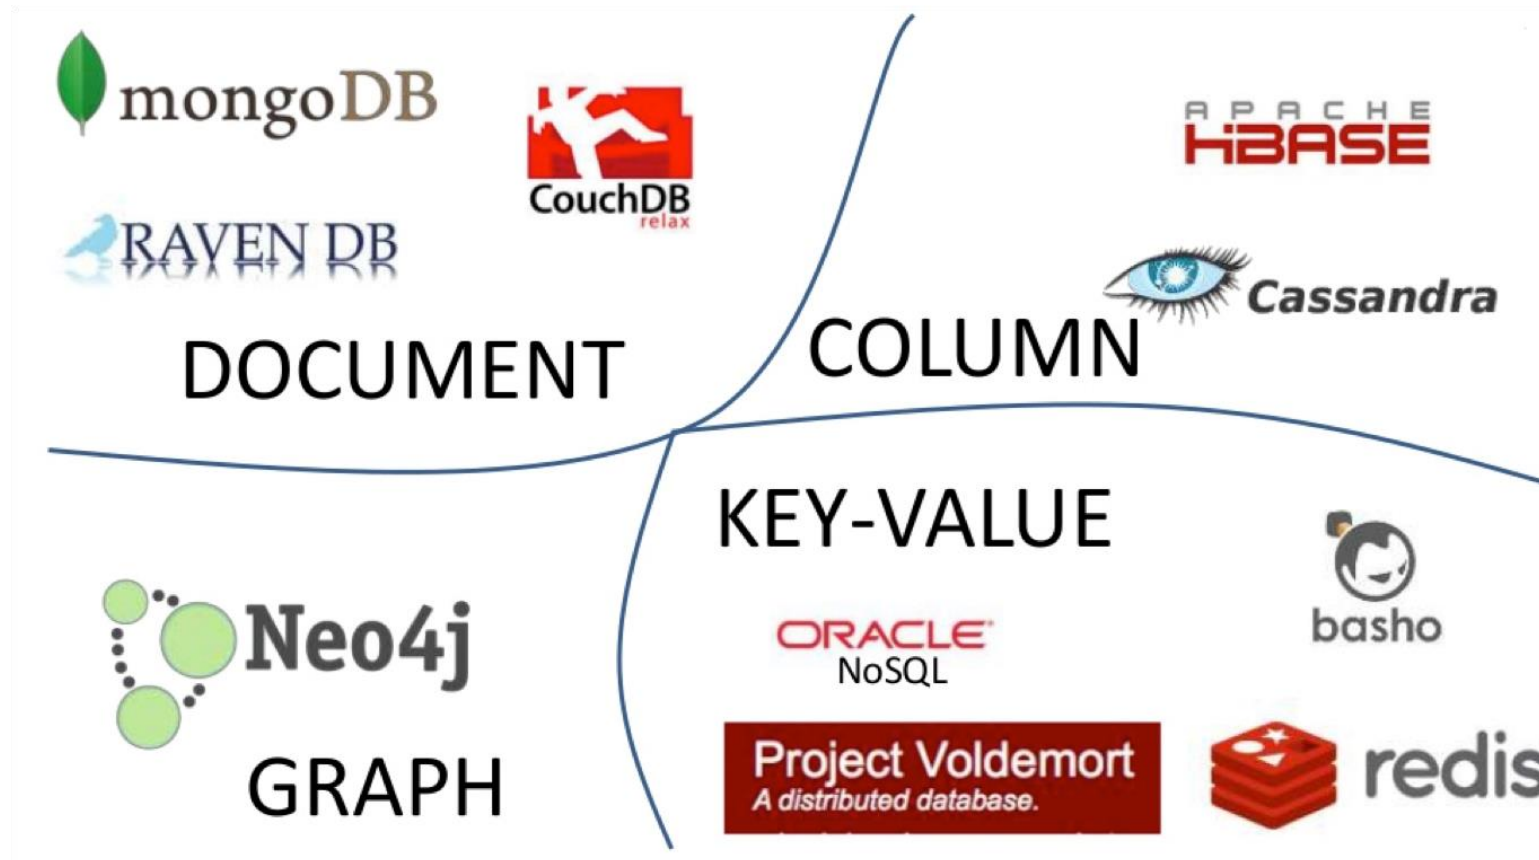
ORACLE

PostgreSQL



SYBASE

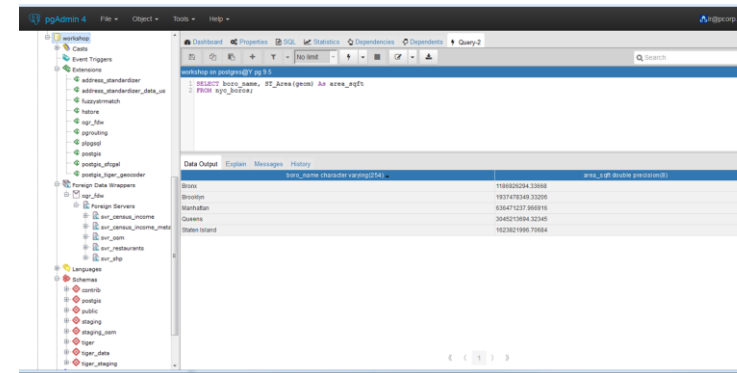




CLI o Graphical client

```
C:\WINDOWS\system32\cmd.exe - redis-server.exe
C:\Tools>redis-server.exe
[7460] 29 May 15:34:54.363 # Warning: no config file specified, using the default config.
[7460] 29 May 15:34:54.363 #
Redis 2.8.4 (00000000/0) 64 bit
Running in stand alone mode
Port: 6379
PID: 7460
http://redis.io

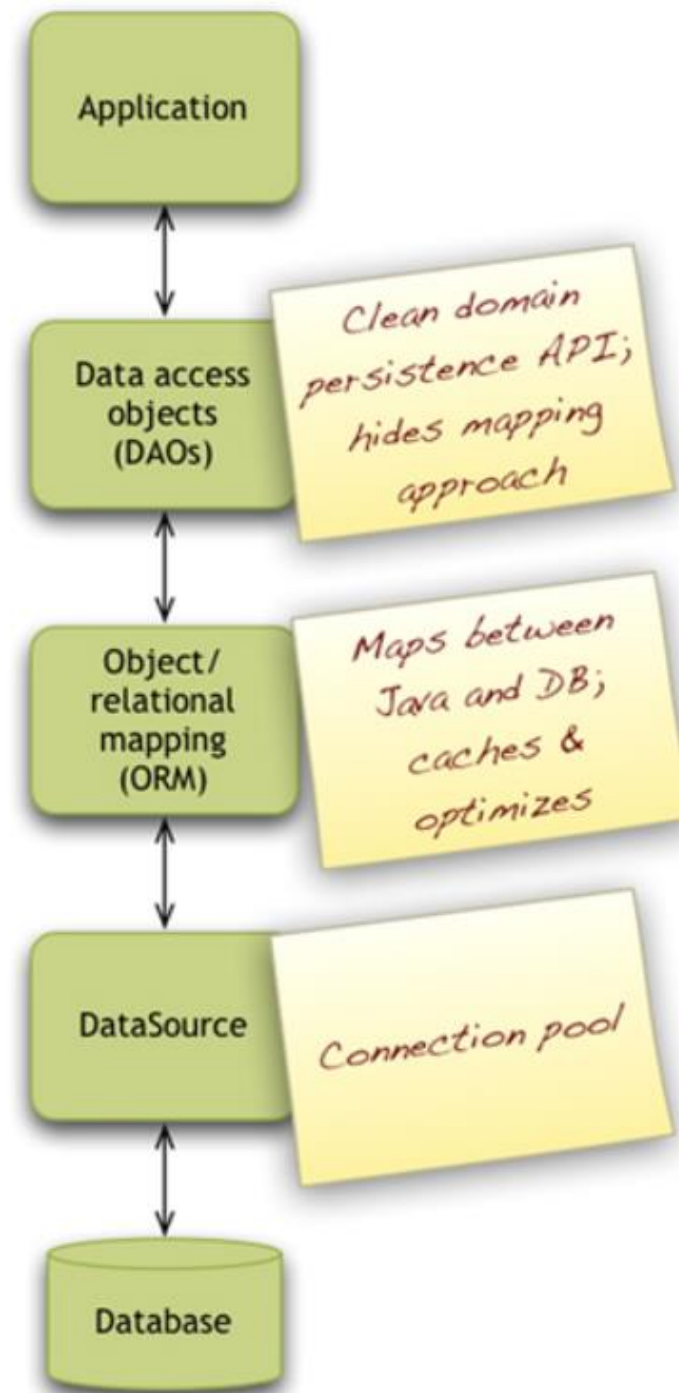
[7460] 29 May 15:34:54.385 # Server started, Redis version 2.8.4
[7460] 29 May 15:34:54.386 # DB loaded from disk: 0.001 seconds
[7460] 29 May 15:34:54.386 # The server is now ready to accept connections on port 6379
```



Robomongo



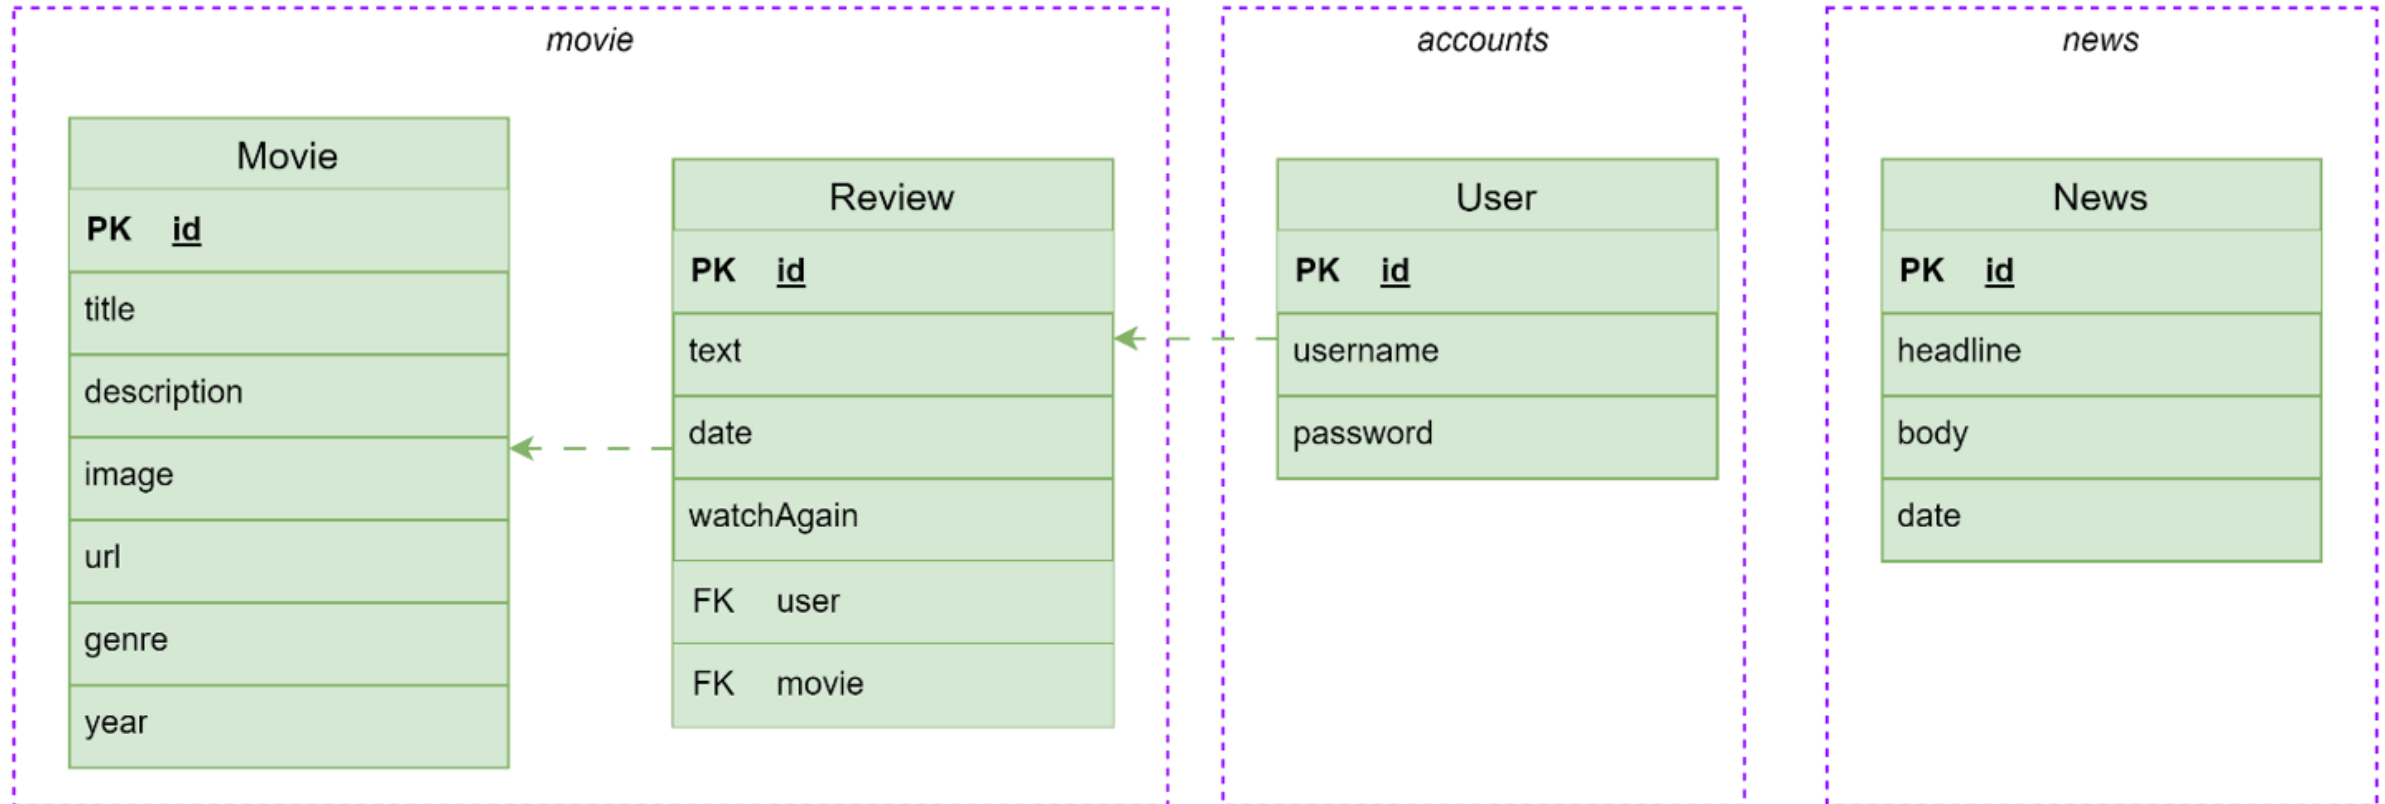
Database Connectivity



ORM (Object-relational mapping)

- System to map any data to the object structure.
 - Frameworks mostly use ORM to interface to user code and covers the problematic to make the storage of objects.
 - The mapping can be defined using XML files or metadata annotations.
-
- JPA (Java Persistence API)
 - ORM with Spring

Example – MovieReviews:



Activity (to report on the wiki):

Refine the domain model to design the Relational Database Model of your project.

Questions???

https://eafit-my.sharepoint.com/:v:/g/personal/pvallej3_eafit_edu_co/EUZnM6VwZclIpHpQSDaaS24BLAbrZ14-giwKLHSwxPYuyA?e=hCiJ6k&nav=eyJyZWZlcnJhbEluZm8iOmsicmVmZXJyYWxBcHAiOiJTdHJlYW1XZWJBcHAiLCJyZWZlcnJhbFZpZXciOiJTdGFyZURpYWxvZy1MaW5rliwicmVmZXJyYWxBcHBQbGF0Zm9ybSI6IldlYiIsInJlZmVycmFsTW9kZSI6InZpZXcifX0%3D